

# Clojure & AWS Lambda: Serverless in the Realm of the Wild Paren

---

# Who?

---

Michael Langford

---

@mj\_langford on twitter

---

michael.langford@rowdylabs.com

---

iOS + Clojure + Creating an App Business

# What?

---

This is a talk about AWS  
Lambda: Its business case  
and its technical use.

---

Lambda is a 64-bit Linux  
compute environment that  
runs JVM, Node or Python  
code.

---

# Designed to be used as "functions"

- Which really are microservices\*
- Especially on the JVM

Can run in 1-2ms  
increments and Amazon  
will bill you for nothing but  
those 1-2ms

---



## Can be Called From:

IoT buttons, Amazon Echo, Web Requests, Things being added to various databases, notifications, queues filling up, the command line, code on client machines, other Lambda functions, etc

# Why?

---

## Your Pain Points:

- Slow code because you don't have enough computers to run it quickly
- Dying servers
- Capacity increasing is manual/hard/slow
- Code that costs thousands per month to "run", when mostly idle

## More Pain Points:

- Sleep you don't get because you're on call if the servers go down
- Angry clients you lose because the servers went down
- Crunch time at work to rearchitect for scale

# An End To These Pains?

## Some of them

- Not enough computers: Solved
- Servers failing: Solved (kinda)
- Too Long to Add Capacity: Solvable
- Cheap to run per month: (kinda, but at what price)

## Some of them

- Sleep due to pager duty, etc: Solved (mostly)
- Angry clients due to failing servers: Solved
- Crunch time at work due to rearchitecting: Solved

# What about cost?

- \$\$\$\$?
- Cold hard cash?



## Kinda?

- You're building for scale from the get-go when you're choosing this different, constrained architecture.
- Slower work at first to save on scaling later.
- Possible loss for any organization continually staffed by software developers and system administrators (at first).

## But in some cases...

- "Continually staffed" is the kicker though.
- Like many "autoscaling" PaaSs, it's a superior tool for companies without an immediate support structure.
- Mobile gets fewer "second chances" when it fails to scale.

## Alternatives that also solve some of these pains

- Older PaaS (there are tons, but Google App Engine and Heroku are two examples)
- Scaling Clouds (even AWS products like ElasticBeanStalk)
- Backend as a Service (Cloudkit, Firebase, Parse)

## **Advantages over those:**

- Less ability to fail in a way that looks bad.
- Dedicated to autoscaling architecture.
- Less vulnerable to a BaaS going out of business (Parse).
- Simpler client software when logic in the cloud (compared to Baas).
- Designed for connection to everything without coding it yourself.

# How?

---

# Some Architectures

(user) <-> (software client) <-> (lambda) <->  
(dynamodb)

(user) <-> (software client) <-> (api gateway) <->  
(lambda) <-> (dynamodb)

## Some More Architectures

(user) <-> (IoT Button) --> (simple queue service) <->  
(lambda) <-> (dynamodb) <-> (analytics user)

(user) <-> (software client) <-> (api gateway) <->  
(lambda) <-> (relational database)

## Still More Architectures

(user) <-> (software client) --> (s3 bucket) --> (lambda)  
--> (another s3 bucket)

(user) --> (software client) --> (simple queue service) --  
> (lambda) --> (simple notification service) -->  
(software client) --> (another person)



# The One We're Doing Tonight

(user) --- (software client) --- (api gateway) --- (lambda)  
--- (dynamodb)

**So let's write a system using this architecture**

Melvin: signs up for an event

Sees himself and others that have signed up for the event

Meets people at the event, doing the activity

# Skipping IAM setup

- Identity Access Management
- Running the commands in this presentation will require a overprivledged user
- Delegate most of these privledges to your CI tool and constrain the dev roles
- <http://docs.aws.amazon.com/lambda/latest/dg/intro-permission-model.html>

# Our story so far

(person:Melvin) --- (software client) --- (api gateway)  
--- (lambda) --- (dynamodb)

# Software Client

---

# Let's do a simple JSON API with curl

(user:Melvin) --- (software client:command line client script) --- (api gateway) --- (lambda) --- (dynamodb)

**Let's talk about the pain points now.**

Wow, Amazon API Gateway hurts. It's hard to configure sensibly, doesn't propagate HTTP errors well, and is just super raw in many places.

## What is it in theory?

- You use a web GUI or swagger specification file to specify a bunch of endpoints, including what http methods work on it
- You deploy it
- You assign other AWS services to handle the requests coming in, and return appropriate status codes



## What is it in practice?

- You use a very hard to use, somewhat misdocumented GUI that is light on examples.
- Oriented towards JSON-style APIs, not HTTP param-type APIs.
- Opaque when it fails.
- The ring adapter requires a wierd pile of JS to paste in there.

## What are the good parts?

- HTTPS by default.
- Able to use custom authorizers or Cognito (Amazon's user authentication system).
- Reasonable to tie into CloudFront.
- Can dispatch to EC2 servers, Lambda, or a mix.
- Generates "okay" iOS/Android/JS API.

## Tricks:

- Configure them in the order specified in `jpb/ring-aws-lambda-adapter`'s readme
- Look at the names in the upper right while selecting a method on a resource
- Copy that large block of JS, because they say to.
- You can't export YAML until you deploy, chose YAML for swagger + api
- Edit the YAML and use `lein aws-api-gateway update`

# The Story So Far

(user:Melvin) --- (software client:Cljs client project) ---  
(api gateway:arctic-fun-times-api) --- (lambda) ---  
(dynamodb)

# Lambda

---

# LAMBDA (finally)

Can be a JVM, Node or Python (2.7) "function"

Think of the "function" as the entry point into a microservice...I'll get back to why it's not a function.

**Do not use:**

gen-class example from amazon/pojo example

**Probably do not use:**

Lambada (the extra 'a' is not a typo)



**Instead: ring adapter lambda (less hard way)**

<https://github.com/jpb/ring-aws-lambda-adapter>

## Why Ring Adapter Lambda:

- You can write what works locally as a web app...and on AWS as a Lambda function.
- Main is entirely for running locally.
- In theory, I guess you could deploy the app to EC2/EBS if you want...or Lambda.

**Amazon puts your code on ice and shuts down the handlers if it ever isn't called enough**

Which isn't a huge problem for node or python...but for JVM langs this is death.

**Not just slow?**

No, death. The API Gateway has a hard timeout of 10 seconds. If you don't return in 10 seconds, it errors out.

**That's actually too slow to start up some Clojure servers  
on Lambda**

Oi. So what's the remedy?

# The Fix:

- After deploy, invoke the function once
- Have your "function" called once a minute (eats 1/3 of your "free" calls)



Configure your Lambda function **backendProduction** to respond to events from the selected trigger. Click on the box below to select your trigger type.

### CloudWatch Events - Schedule

**Rule name** ping\_me ⓘ

**Rule description** pings to keep hot ⓘ

**Schedule expression** rate(1 minute) ⓘ



### Lambda

Lambda will add the necessary permissions for CloudWatch Events to invoke your Lambda function on a schedule. [Learn more](#) about the Lambda permissions model.

**Enable trigger** ☒ ⓘ

# The Story So Far

(user:Melvin) --- (software client:command line client script) --- (api gateway:arctic-fun-times-api) --- (lambda:backendProduction) --- (dynamodb)

# DynamoDB



# DynamoDB

- Lambda sometimes drowns traditional relational DBs with connection attempts.
- DynamoDB is not necessary, nor always desired, but it works for many projects.
- Limited cross table queries, weird primary keys, among other oddities.

## The Process:

Create a table, specifying how many concurrent readers and writers we want (to pay for), as well as the weird primary key (plus any secondary "keys" we need to query on).

# Client Opts: Where the DB is + login keys

```
7 ;You generally want to never put keys in source files. Use profile.clj or
8 ; a manager like awspm. If you use a profile.clj, do your team a favor
9 ; and include a profile.clj.template file in source control
10 (def client-opts
11   {:access-key (:ddb-access env)
12    :secret-key (:ddb-secret env)
13    :endpoint    (:ddb-endpoint env)}) ; endpoint is datacenter/region specific
```

# DynamoDB Table Creation Minutia

```
21
22 (def event-table
23   {:pk [:event :s]
24    :opts { ;range-keydef is used to make a secondary part of the query key
25            ; you can only query on the primary key or range-keydef keys
26            ; consider using a relational database if you need more complex
27            ; queries
28            :range-keydef [:name :s]
29            :block? true
30            :throughput {:read 5
31                          :write 5}}})
32
```

# Ensure the table is there on every execution to pick up changes

```
33 (def table-defintions
34   {:Events event-table})
35
36 (defn setup
37   "this makes sure all the tables are there we use elsewhere"
38   []
39   (println "For DB in location " (:endpoint client-opts)))
40   (doseq [[table-name t] table-defintions]
41     (far/ensure-table client-opts
42                       table-name
43                       (:pk t)
44                       (:opts t))))
45
46
```

# The Story: Fin

(user:Melvin) --- (software client:command line client script) --- (api gateway:arctic-fun-times-api) --- (lambda:backendProduction) --- (dynamodb:Events)

# Neglected Topics

- IAM
- Testing
- Production/Staging Separation
- Local DynamoDB
- CloudWatch
- User Management
- JSON Client App

# Demo

Code: <https://github.com/langford/clj-aws-lambda-example>



# Thanks!

---

References: See the code for all the libraries, clojars/  
github for lambada

Michael Langford

---

@mj\_langford on twitter

---

michael.langford@rowdylabs.com

---

iOS + Clojure + Creating an App Business

